

Rvalue References, Forwarding, Move Semantics

4/24/2005

1

A Generic Factory Function

```
template <class T, class A1>
std::auto_ptr<T> factory(A1 const& x1)
{
    return std::auto_ptr<T>( new T(x1) );
}

class widget
{
public:
    widget(int size);
    ...
};
// yay!
std::auto_ptr<widget> w = factory<widget>(10);
```

4/24/2005

copyright 2005 David Abrahams

2

First Problem

```
class traced : boost::noncopyable
{
public:
    traced(std::ostream& log) : log(log)
    {
        log << "constructed" << std::endl;
    }
    ~traced()
    {
        log << "destroyed" << std::endl;
    }
    // ...
private:
    ostream& log;
};
// boo!
std::auto_ptr<traced> m = factory<traced>(std::cout);
```

4/24/2005

copyright 2005 David Abrahams

3

Fix: Add an Overload

```
template <class T, class A1>
std::auto_ptr<T> factory(A1& x1)
{
    return std::auto_ptr<T>( new T(x1) );
}

template <class T, class A1>
std::auto_ptr<T> factory(A1 const& x1)
{
    return std::auto_ptr<T>( new T(x1) );
}

// yay!
std::auto_ptr<traced> m = factory<traced>(std::cout);
std::auto_ptr<widget> w = factory<widget>(10);
```

4/24/2005

copyright 2005 David Abrahams

4

First Problem

```
class traced : boost::noncopyable
{
public:
    traced(std::ostream& log) : log(log)
    {
        log << "constructed" << std::endl;
    }
    ~traced()
    {
        log << "destroyed" << std::endl;
    }
    // ...
private:
    ostream& log;
};
// boo!
std::auto_ptr<traced> m = factory<traced>(std::cout);
```

4/24/2005

copyright 2005 David Abrahams

5

Handling Two Arguments

```
template <
    class T, class A1, class A2
>
std::auto_ptr<T>
factory(A1& x1, A2& x2)
{
    return std::auto_ptr<T>(
        new T(x1, x2) );
}

template <
    class T, class A1, class A2
>
std::auto_ptr<T>
factory(A1 const& x1, A2& x2)
{
    return std::auto_ptr<T>(
        new T(x1, x2) );
}
```

```
template <
    class T, class A1, class A2
>
std::auto_ptr<T>
factory(A1& x1, A2 const& x2)
{
    return std::auto_ptr<T>(
        new T(x1, x2) );
}

template <
    class T, class A1, class A2
>
std::auto_ptr<T>
factory(A1 const& x1, A2 const& x2)
{
    return std::auto_ptr<T>(
        new T(x1, x2) );
}
```

4/24/2005

copyright 2005 David Abrahams

6

Handling N Arguments

Uh oh...

Handling N arguments requires 2^N overloads!

N Overloads for N Args (Almost)

```
template <class T, class A1>
std::auto_ptr<T> factory(A1& x1)
{
    return std::auto_ptr<T>( new T(x1) );
}

// yay!
std::auto_ptr<traced> m = factory<traced>(std::cout);

// yay!
int const x = 10;
std::auto_ptr<wi dget> w = factory<wi dget>(x);

// boo!
std::auto_ptr<wi dget> w = factory<wi dget>(10);
```

What's The Problem?

- 8.5.3 paragraph 5 (paraphrasing):
A non-const reference shall not be bound to an rvalue
- Quoting non-normative text
[Example:
double& rd2 = 2.0; // error: not an lvalue and reference not const
]
- There is a gap in the type system
(expressiveness, not safety).

3.10 Lvalues and Rvalues

1. Every expression is either an lvalue or an rvalue.
2. An lvalue refers to an object or function. Some rvalue expressions – those of class or cv-qualified class type – also refer to objects.
3. [Note: some built-in operators and function calls yield lvalues. [Example: if E is an expression of pointer type, then *E is an lvalue expression referring to the object or function to which E points. As another example, the function
int& f();
yields an lvalue, so the call f() is an lvalue expression.]]
4. [Note: some built-in operators expect lvalue operands. [Example: built-in assignment operators all expect their left hand operands to be lvalues.]...]

14.3 [Note: a reference can be thought of as a name of an object.]

Lvalues and Rvalues (Summary)

- All expressions that are names or that are references are lvalues.
- Everything else is an rvalue

```
int x;  
int f();  
int& g();
```

```
x;    // L  
10;   // R  
f();  // R  
g();  // L
```

4/24/2005

copyright 2005 David Abrahams

11

Why This Rule? – Killer Example

B. Stroustrup, *The **Design and Evolution** of C++* ("D&E"), section 3.7

```
void incr(int& rr) { rr++; }
```

```
void g()  
{  
    double ss = 1;  
    incr(ss);  
}
```

4/24/2005

copyright 2005 David Abrahams

12

We Need an Über-Reference!

- The rvalue reference fills that role
- Spelled: X&&
- Binds to both lvalues and rvalues

```
template <class T, class A1, class A2>
std::auto_ptr<T> factory(A1&& x1, A2&& x2)
{
    return std::auto_ptr<T>(new T(x1, x2));
}
```

- This code isn't quite right yet (more later)

Move Semantics

Sometimes you don't need to copy.

– B. Stroustrup

Simple Examples

```
std::string to_string(int x)
{
    std::string x
        = boost::lexical_cast<std::string>(x);
    return x;
}

char first(std::string x)           // copy
{
    std::sort(x.begin(), x.end());
    return *x.begin();
}

std::string greet("hello, world!");
char y = first(greet);

char z = first(to_string(10));
```

4/24/2005

copyright 2005 David Abrahams

15

Worse Examples

```
std::vector<std::string> explode(std::string const& x)
{
    std::vector<std::string> result;

    std::transform(
        x.begin(), x.end(), std::back_inserter(result),
        to_string);

    return result;
}

void insert_in_order( // Precondition: v is sorted
    std::vector<std::string>& v, std::string const& s)
{
    v.insert(lower_bound(v.begin(), v.end(), s), s);
}
```

4/24/2005

copyright 2005 David Abrahams

16

Making std::string Moveable

```
class string {
public:
    // copy semantics
    string(const string& s)
        : data(new char[s.size]), size (s.size)
    { memcpy(data, s.data, size); }

    string& operator=( const string& s );
    {
        // move constructor
        string(string&& s)
            : data_(s.data), size_(s.size)
        { s.data = 0; s.size = 0; }

        // move assignment
        string& operator=(string&& s)
        { swap(s); return *this; }

private:
    char* data; size_t size;
};
```

4/24/2005

copyright 2005 David Abrahams

17

Making std::vector move-aware

```
// This is the usual insert (no moving)
iterator insert(iterator pos, value_type const& x)
{
    if (capacity() > size())
    {
        if (pos == end())
            construct(x, end());
        else
        {
            construct(*end() - 1, end());
            copy_backward(pos, end() - 1, pos + 1);
            *pos = x;
        }
    }
    else // ...reallocate and copy the buffer...
}
```

4/24/2005

copyright 2005 David Abrahams

18

1. Handle the “x is an Rvalue” Case

```
// This is an overload
iterator insert(iterator pos, value_type&& x)
{
    if (capacity() > size())
    {
        if (pos == end())
            construct(x, end());
        else
        {
            construct(*end() - 1, end());
            copy_backward(pos, end() - 1, pos + 1);
            *pos = x;
        }
    }
    else // ...reallocate and copy the buffer...
}
```

4/24/2005

copyright 2005 David Abrahams

19

2. Explicit Moving in construct

```
// The old way
template <class T>
void construct(T const& x, void* p)
{ new (p) T(x); }

// Overload for move
template <class T>
void construct(T&& x, void* p)
{ new (p) T( std::move(x) ); }

// “Cast to rvalue – I know it’s safe to move”
template <class T>
typename remove_reference<T>::type&&
move(T&& t)
{
    return t;
}
```

4/24/2005

copyright 2005 David Abrahams

20

3. Explicit Moving in vector

```
// Do this to both overloads
iterator insert(iterator pos, value_type&& x)
{
    if (capacity() > size())
    {
        if (pos == end())
            construct(x, end());
        else
        {
            construct(std::move(*end()-1), end());
            move_backward(pos, end() - 1, pos + 1);
            *pos = x;
        }
    }
    else // ...reallocate and move the buffer...
}
```

4/24/2005

copyright 2005 David Abrahams

21

move_ptr / unique_ptr

- Why don't we allow std::auto_ptr in standard containers?
- Why is move_ptr okay in standard containers?
- Why might this be better than using shared_ptr?

4/24/2005

copyright 2005 David Abrahams

22

Move is a Pure Optimization

- Making types movable is entirely optional
- Library can be transparently upgraded to use move
- No performance decrease for existing code
- Probably a performance boost, even if you don't make your types movable.
- Libraries use "move if available, else copy" strategy: no moveability detection required.
- A class with move assign but no move construct (or vice-versa) is just fine

4/24/2005

copyright 2005 David Abrahams

23

Arbitrarily Huge Performance Boost

- Just make a sufficiently gnarly data structure:

```
std: : map<  
    std: : set<std: : string>,  
    std: : vector<int>  
>
```

4/24/2005

copyright 2005 David Abrahams

24

Moving and Forwarding

- Our last factory isn't perfect:

```
template <class T, class A1, class A2>
std::auto_ptr<T> factory(A1&& x1, A2&& x2)
{
    return std::auto_ptr<T>(new T(x1, x2));
}

struct Foo
{
    foo(std::string x, std::vector<int> y);
};

factory<Foo>(
    std::string("Movel t"), std::vector<int>(10, 0));
```

4/24/2005

copyright 2005 David Abrahams

25

Perfect Forwarding

- Need a way to know and restore rvalue-ness of argument
- Rvalue-ness captured in A1 and A2

```
template <class T, class A1, class A2>
std::auto_ptr<T> factory(A1&& x1, A2&& x2)
{
    return std::auto_ptr<T>(
        new T(
            std::forward<A1>(x1),
            std::forward<A2>(x2)
        )
    );
}
```

4/24/2005

copyright 2005 David Abrahams

26

Extended Reference-Collapsing Rules

- `A& &` → `A&` by CWG 106
- `A& &&` → `A&` for perfect forwarding
- `A&& &` → `A&` for completeness
- `A&& &&` → `A&&` for perfect forwarding

Thank you, Peter Dimov!

Definition of `std::forward`

```
template <class T>
struct unspecified
{
    typedef T type;
};

template <class T>
T&& forward(typename unspecified<T>::type&& t)
{
    return t;
}
```

Library-Only Move Solutions

- History:
 - `std::auto_ptr`
 - MOJO (Alexajdrescu)
 - Proposed Boost.Move (Abrahams)
- Close, but no cigar:
 - Writing movable classes is arcane and verbose (conversion operator, two copies of lvalue copy/assign)
 - Ugly: tricks on the type system to do something the compiler already knows about
 - All optimizations lost at forwarding boundaries (`tr1::bind`, `factory`, `bind1st`, `bind2nd`, `vector::push_back`)

4/24/2005

copyright 2005 David Abrahams

29

Move and Inheritance

```
struct Base
{
    Base(Base&& b);
};

struct Derived
    : Base
{
    Derived(Derived&& d)
        : Base(std::move(d)) {}
};
```

4/24/2005

copyright 2005 David Abrahams

30